

Importing Data into Python from Relational Databases



Xavier Morera

BIG DATA INC.

@xmorera www.xaviermorera.com



Importing sqlite Database Files



Database engine

- Store & work with relational data
- Simple to use and portable

Library and the db file

- Use sqlite3 in Python

Many notable users

```
import sqlite3
# ls
stack_connection = sqlite3.connect('importing_sqlite.db')
type(stack_connection)
stack_cursor = stack_connection.cursor()
stack_cursor.execute("select name from sqlite_master where type = 'table';")
stack_cursor.fetchone()
```

Importing sqlite Database Files



Import **sqlite3** and create a connection using **connect**

Create a **Cursor** object, and **execute**

Retrieve results using **fetchone** or **fetchall**




```
# ls -l
```

```
stack_connection_bad = sqlite3.connect('bad_name_sqlite.db')  
stack_connection_bad.cursor().execute("select name from sqlite_master where type =  
'table';").fetchall()
```

```
# ls -l
```

Watch Out!



Make a mistake on the database name

And a new, empty database may be created



```
rows = stack_cursor.execute('select * from posts').fetchall()
type(rows)
rows[0]
type(rows[0])
stack_cursor.execute('select * from posts limit 1').fetchall()
stack_cursor.execute('select Id, Score, Tags from posts limit 3').fetchall()
```

Querying Data



Explore your data

- Returns a **list** of objects of type **tuple**

Refine your queries



Table of Contents

sqlite3 — DB-API 2.0
interface for SQLite databases

- Module functions and constants
- Connection Objects
- Cursor Objects
- Row Objects
- Exceptions
- SQLite and Python types
 - Introduction
 - Using adapters to store additional Python types in SQLite databases
 - Letting your object adapt itself
 - Registering an adapter callable
 - Converting SQLite values to custom Python types
 - Default adapters and converters
- Controlling Transactions
- Using sqlite3 efficiently
 - Using shortcut methods
 - Accessing columns by name instead of by index
 - Using the connection as a context manager
- Common issues

sqlite3 — DB-API 2.0 interface for SQLite databases

Source code: [Lib/sqlite3/](#)

SQLite is a C library that provides a lightweight disk-based database that doesn't require a separate server process and allows accessing the database using a nonstandard variant of the SQL query language. Some applications can use SQLite for internal data storage. It's also possible to prototype an application using SQLite and then port the code to a larger database such as PostgreSQL or Oracle.

The sqlite3 module was written by Gerhard Häring. It provides a SQL interface compliant with the DB-API 2.0 specification described by [PEP 249](#).

To use the module, you must first create a [Connection](#) object that represents the database. Here the data will be stored in the `example.db` file:

```
import sqlite3
conn = sqlite3.connect('example.db')
```

You can also supply the special name `:memory:` to create a database in RAM.

Once you have a [Connection](#), you can create a [Cursor](#) object and call its `execute()` method to perform SQL commands:

```
c = conn.cursor()

# Create table
c.execute('''CREATE TABLE stocks
            (date text, trans text, symbol text, qty real, price real)''')

# Insert a row of data
c.execute("INSERT INTO stocks VALUES ('2006-01-05', 'BUY', 'RHAT', 100, 35.14)")

# Save (commit) the changes
```

```
import pandas as pd

stack_connection = sqlite3.connect('importing_sqlite.db')

posts_df = pd.read_sql("select * from posts;", stack_connection)

type(posts_df)

posts_df.columns

posts_df.head()
```

Taking Advantage of Pandas

Working with a list of tuples may not be ideal

Enter **pandas**

- Still with **sqlite3**, but you create a **DataFrame**
- Use **read_sql**



Table of Contents

sqlite3 — DB-API 2.0
interface for SQLite databases

- Module functions and constants
- Connection Objects
- Cursor Objects
- Row Objects
- Exceptions
- SQLite and Python types
 - Introduction
 - Using adapters to store additional Python types in SQLite

sqlite3 — DB-API 2.0 interface for SQLite databases

Source code: [Lib/sqlite3/](#)

SQLite is a C library that provides a lightweight disk-based database that doesn't require a separate server process and allows accessing the database using a nonstandard variant of the SQL query language. Some applications can use SQLite for internal data storage. It's also possible to prototype an application using SQLite and then port the code to a larger database such as PostgreSQL or Oracle.

The sqlite3 module was written by Gerhard Häring. It provides a SQL interface compliant with the DB-API 2.0 specification described by [PEP 249](#).

[pandas 0.23.4 documentation](#) » [API Reference](#) »

[previous](#) | [next](#) | [modules](#) | [index](#)

Table Of Contents

What's New
Installation
Contributing to pandas
Package overview
10 Minutes to pandas
Tutorials
Cookbook
Intro to Data Structures
Essential Basic Functionality
Working with Text Data
Options and Settings
Indexing and Selecting Data
MultiIndex / Advanced Indexing
Computational tools
Working with missing data
Group By: split-apply-combine
Merge, join, and concatenate

pandas.DataFrame

`class pandas.DataFrame(data=None, index=None, columns=None, dtype=None, copy=False)` [\[source\]](#)

Two-dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns). Arithmetic operations align on both row and column labels. Can be thought of as a dict-like container for Series objects. The primary pandas data structure.

Parameters:

data : *numpy ndarray (structured or homogeneous), dict, or DataFrame*
Dict can contain Series, arrays, constants, or list-like objects
Changed in version 0.23.0: If data is a dict, argument order is maintained for Python 3.6 and later.

index : *Index or array-like*
Index to use for resulting frame. Will default to RangeIndex if no indexing information part of input data and no index provided

columns : *Index or array-like*

Working with Databases Using SQLAlchemy



The Python SQL toolkit

Supports

- SQLite, PostgreSQL, MySQL
- Oracle, MS SQL, Firebird, Sybase...



```
from sqlalchemy import create_engine
engine = create_engine('sqlite:///importing_sqlite.db')
type(engine)
dir(engine)
engine.table_names()
engine.url
engine.dialect
engine.driver
```

Using SQLAlchemy with SQLite



Start with the necessary import

And then `create_engine`

- Provide connection string
- You have an engine



Anatomy of a Connection String

```
engine = create_engine('sqlite:///importing_sqlite.db')
```



Anatomy of a Connection String

```
engine = create_engine('sqlite:///importing_sqlite.db')
```



Anatomy of a Connection String

`sqlite:///importing_sqlite.db`



Anatomy of a Connection String

`dialect:///importing_sqlite.db`



Anatomy of a Connection String

`dialect:///dbname`



Anatomy of a Connection String

`dialect[+driver]://user:password@hostname/dbname`



Anatomy of a Connection String

`dialect[+driver]://user:password@hostname/dbname[?key=value]`



```
engine_sqlite = create_engine('sqlite:///importing_sqlite.db')  
  
engine_mysql = create_engine('postgresql://xavier:postgres@localhost:5432/importing_postgres')  
  
engine_postgresql = create_engine('mysql+mysqlconnector://root:mysql@localhost:3306/importing_mysql')
```

Connecting to Your Database of Choice

Layer of abstraction

Create database agnostic applications

Load into a Pandas DataFrame



```
engine_sqlite = create_engine('sqlite:///importing_sqlite.db')
```

```
engine_postgres = create_engine('postgresql://xavier:postgres@localhost:5432/importing_postgres')
```

```
engine_mysql = create_engine('mysql+mysqlconnector://root:mysql@localhost:3306/importing_mysql')
```





Open source database

Runs on virtually all major platforms

- Top 3 of most widely used

Client-server model



```
# show databases;  
# use importing_mysql  
# show tables;  
engine = create_engine('mysql+mysqlconnector://root:mysql@localhost:3306/importing_mysql')  
posts = pd.read_sql_table('posts', engine, index_col='Id')  
type(posts)  
posts.columns  
posts.head()
```

Importing Data with Pandas



Use **SQLAlchemy** to connect to MySQL with pandas

- Create engine
- Requires **mysql-connector-python**

Load table into **DataFrame** with **read_sql_table** and set index column



Table Of Contents

What's New
Installation
Contributing to pandas
Package overview
10 Minutes to pandas
Tutorials
Cookbook
Intro to Data Structures
Essential Basic Functionality
Working with Text Data
Options and Settings
Indexing and Selecting Data
MultiIndex / Advanced Indexing
Computational tools
Working with missing data
Group By: split-apply-combine
Merge, join, and concatenate
Reshaping and Pivot Tables
Time Series / Date functionality
Time Deltas
Categorical Data
Visualization
Styling
IO Tools (Text, CSV, HDF5, ...)

- CSV & Text files
- JSON
- HTML
- Excel files
- Clipboard
- Pickling
- msgpack
- HDF5 (PyTables)
- Feather
- Parquet
- SQL Queries

pandas.read_sql_table

`pandas.read_sql_table(table_name, con, schema=None, index_col=None, coerce_float=True, parse_dates=None, columns=None, chunksize=None)` [\[source\]](#)

Read SQL database table into a DataFrame.

Given a table name and a SQLAlchemy connectable, returns a DataFrame. This function does not support DBAPI connections.

Parameters:

table_name : *string*

Name of SQL table in database.

con : *SQLAlchemy connectable (or database string URI)*

SQLite DBAPI connection mode not supported.

schema : *string, default None*

Name of SQL schema in database to query (if database flavor supports this). Uses default schema if None (default).

index_col : *string or list of strings, optional, default: None*

Column(s) to set as index(MultiIndex).

coerce_float : *boolean, default True*

Attempts to convert values of non-string, non-numeric objects (like decimal.Decimal) to floating point. Can result in loss of Precision.

parse_dates : *list or dict, default: None*

- List of column names to parse as dates.
- Dict of {column_name: format string} where format string is strftime compatible in case of parsing string times or is one of (D, s, ns, ms, us) in case of parsing integer timestamps.
- Dict of {column_name: arg dict}, where the arg dict corresponds to the keyword arguments of `pandas.to_datetime()` Especially useful with databases without native Datetime support, such as SQLite.

columns : *list, default: None*

```
posts = pd.read_sql_table('posts', engine, columns=['Id', 'CreationDate', 'Tags'])
posts.head()
type(posts.iloc(1)[1])
posts = pd.read_sql_table('posts', engine, columns=['Id', 'CreationDate', 'Tags'],
parse_dates={'CreationDate': {'format': '%Y-%m-%dT%H:%M:%S.%f'}})
type(posts.iloc(1)[1])
```

Importing Data with Pandas



Additional functionality

- Select only specific **columns**
- Use **parse_dates**



Table Of Contents

What's New
Installation
Contributing to pandas
Package overview
10 Minutes to pandas
Tutorials
Cookbook
Intro to Data Structures
Essential Basic Functionality
Working with Text Data
Options and Settings
Indexing and Selecting Data
MultiIndex / Advanced Indexing
Computational tools
Working with missing data
Group By: split-apply-combine
Merge, join, and concatenate
Reshaping and Pivot Tables
Time Series / Date functionality
Time Deltas
Categorical Data
Visualization
Styling
IO Tools (Text, CSV, HDF5, ...)

- CSV & Text files
- JSON
- HTML
- Excel files
- Clipboard
- Pickling
- msgpack
- HDF5 (PyTables)
- Feather
- Parquet
- SQL Queries

pandas.read_sql_table

`pandas.read_sql_table(table_name, con, schema=None, index_col=None, coerce_float=True, parse_dates=None, columns=None, chunksize=None)` [\[source\]](#)

Read SQL database table into a DataFrame.

Given a table name and a SQLAlchemy connectable, returns a DataFrame. This function does not support DBAPI connections.

Parameters:

table_name : *string*

Name of SQL table in database.

con : *SQLAlchemy connectable (or database string URI)*

SQLite DBAPI connection mode not supported.

schema : *string, default None*

Name of SQL schema in database to query (if database flavor supports this). Uses default schema if None (default).

index_col : *string or list of strings, optional, default: None*

Column(s) to set as index(MultiIndex).

coerce_float : *boolean, default True*

Attempts to convert values of non-string, non-numeric objects (like decimal.Decimal) to floating point. Can result in loss of Precision.

parse_dates : *list or dict, default: None*

- List of column names to parse as dates.
- Dict of {column_name: format string} where format string is strftime compatible in case of parsing string times or is one of (D, s, ns, ms, us) in case of parsing integer timestamps.
- Dict of {column_name: arg dict}, where the arg dict corresponds to the keyword arguments of `pandas.to_datetime()` Especially useful with databases without native Datetime support, such as SQLite.

columns : *list, default: None*



+

Psycopg2

PostgreSQL is a solid database

Psycopg2

- Popular database adapter
- PostgreSQL + Python



```
# psql
# \l
# \c importing_postgres
# \d
# SELECT "Id", "Title" FROM posts LIMIT 5;
import psycopg2
stack_connection = psycopg2.connect("dbname=importing_postgres user=xavier
host=localhost")
so_cursor = stack_connection.cursor()
```

Importing Data with Psycopg2

Start by importing **psycopg2**

Create a connection using **connect**

Create a **cursor**



```
so_cursor.execute("select * from posts")
first_row = so_cursor.fetchone()
first_row
type(first_row)
rows = so_cursor.fetchall()
rows
type(rows)
```

Importing Data with Psycopg2

Query with **execute** using the **cursor**

Get your results

- Use **fetchone** or **fetchall**



```
stack_connection.commit()  
stack_connection.close()
```

Importing Data with Psycopg2

Remember to **commit** when not using the connection

And **close** when connection no longer needed



```
engine_mysql.table_names()
```

```
nicer_query = "SELECT posts.Id, Users.DisplayName, posts.AnswerCount,  
posts.ViewCount FROM posts INNER JOIN users on  
posts.OwnerUserId=Users.Id ORDER BY posts.ViewCount DESC LIMIT 5;"
```

```
posts = pd.read_sql(nicer_query, engine_mysql)
```

Importing Relational Data

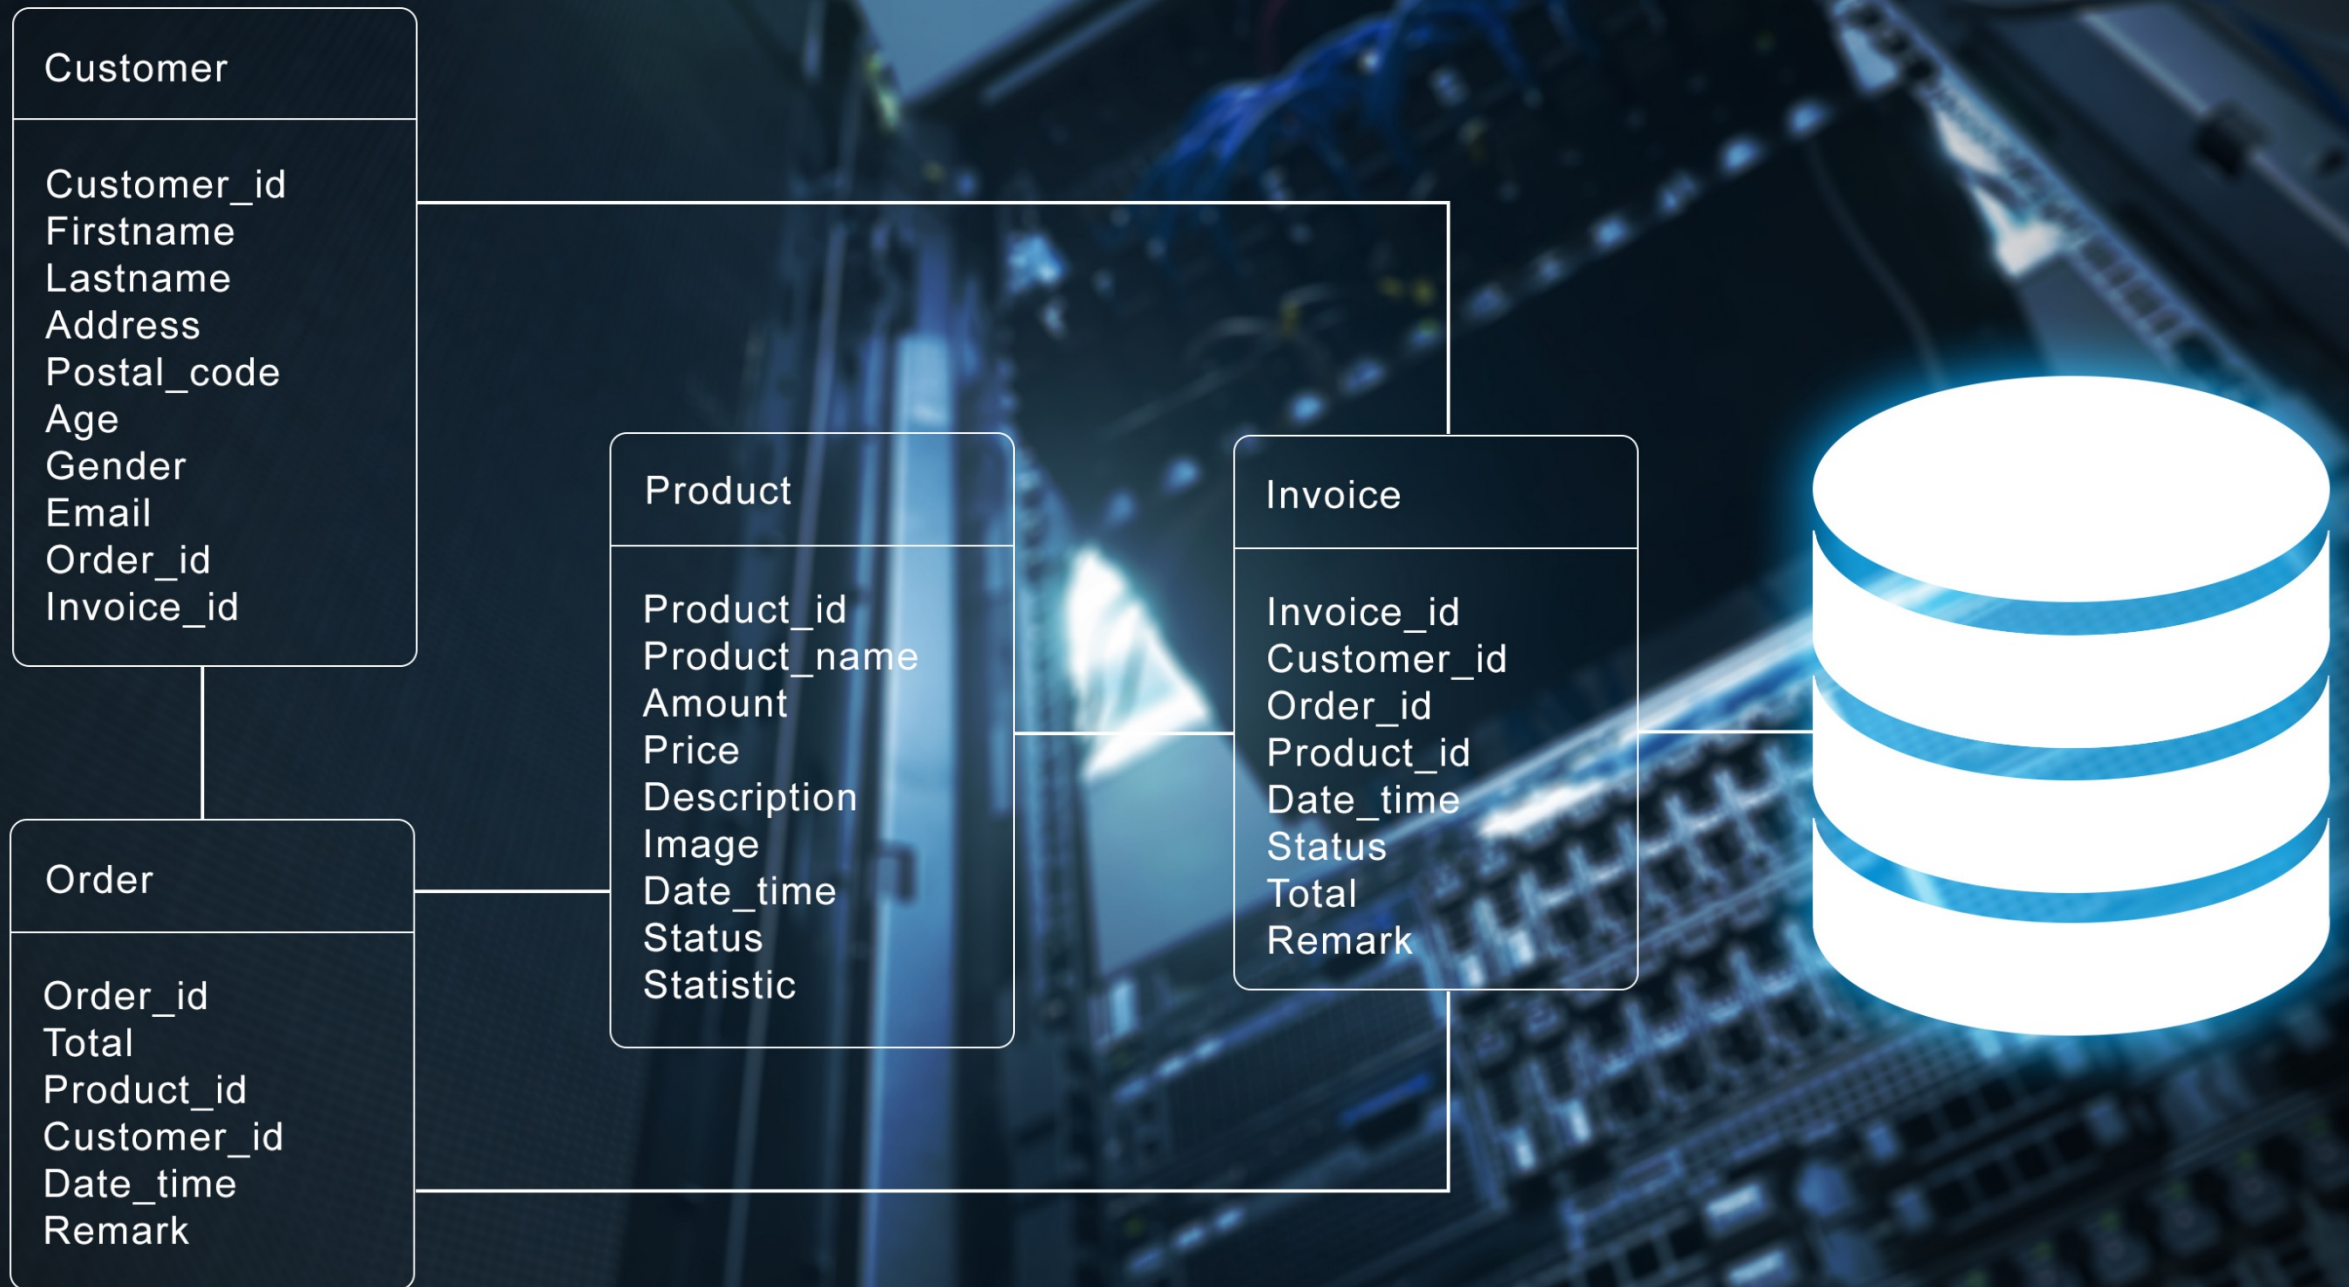
SQL available

- Use it to its full potential

Complex queries, and load into DataFrame

- With `read_sql`





Importing Data: Python Data Playbook



Xavier Morera

BIG DATA INC.

@xmorera www.xaviermorera.com

